

Logistic Regression with R

Xinyu_Chen

2017.2.21

Logistic Regression

Load Data

```
data <- read.table('ex2data1.txt',sep = ',')
X <- data[,c(1,2)]; y <- data[,3]
X <- as.matrix(X)
```

Part 1: Plotting

```
plotData <-
function (X, y, axLabels = c("Exam 1 score","Exam 2 score"), legLabels =
          c('Admitted', 'Not admitted')) {
  #PLOTDATA Plots the data points X and y into a new device
  # PLOTDATA(x,y) plots the data points with + for the positive examples
  # and o for the negative examples. X is assumed to be a Mx2 matrix.

  symbolss <- c(3,21) #plus and empty circle character codes

  ##### This part is for legend to be plotted above the plot
  plot(X[,1],X[,2],type = "n", xaxt = "n", yaxt = "n")
  leg <- legend(
    "topright",legLabels, pch = rev(symbolss),
    pt.bg = "yellow", plot = FALSE
  )

  #custom ylim. Add the height of legend to upper bound of the range
  yrange <- range(X[,2])
  yrange[2] <- 1.04 * (yrange[2] + leg$rect$h)
  #####

  yfac <- factor(y)
  plot(
    X[,1],X[,2], pch = symbolss[yfac] ,bg = "yellow", lwd = 1.3,
    xlab = axLabels[1], ylab = axLabels[2],
    ylim = yrange
  )

  legend("topright",legLabels,pch = rev(symbolss),
        pt.bg = "yellow")
}

cat(sprintf('Plotting data with + indicating (y = 1) examples and o indicating (y = 0) examples.\n'))


## Plotting data with + indicating (y = 1) examples and o indicating (y = 0) examples.
```

```
plotData(X, y)
```

$X[:, 2]$



$X[:, 1]$



Part 2: Compute Cost and Gradient

```
costFunction <- function(X, y) {
  #COSTFUNCTION Compute cost for logistic regression
  # J <- COSTFUNCTION(theta, X, y) computes the cost of using theta as the
  # parameter for logistic regression.

  function(theta) {
    # Initialize some useful values
    m <- length(y) # number of training examples

    J <- 0

    h <- sigmoid(X %*% theta)
    J <- (t(-y) %*% log(h) - t(1 - y) %*% log(1 - h)) / m
    J
  }
}

sigmoid <- function(z) {
  #SIGMOID Compute sigmoid function
  # J <- SIGMOID(z) computes the sigmoid of z.

  z <- as.matrix(z)
  g <- matrix(0, dim(z)[1], dim(z)[2])
```

```

    g <- 1 / (1 + exp(-1 * z))
  }
}

grad <- function(X, y) {
  #COSTFUNCTION Compute gradient for logistic regression
  # J <- COSTFUNCTION(theta, X, y) computes the gradient of the cost
  # w.r.t. to the parameters.
  function(theta) {

    # You need to return the following variable correctly
    grad <- matrix(0,dim(as.matrix(theta)))
    m <- length(y)

    h <- sigmoid(X %*% theta)
    # calculate grads
    grad <- (t(X) %*% (h - y)) / m

    grad
  }
}

# Setup the data matrix appropriately, and add ones for the intercept term
m <- dim(X)[1]
n <- dim(X)[2]

# Add intercept term to x and X_test
X <- cbind(rep(1,m),X)
# Initialize fitting parameters
initial_theta <- rep(0,n+1)

# Compute and display initial cost and gradient
cF <- costFunction(X, y)(initial_theta)
cost <- costFunction(X, y)(initial_theta)
grd <- grad(X,y)(initial_theta)

cat(sprintf('Cost at initial theta (zeros): %f\n', cost))

## Cost at initial theta (zeros): 0.693147

cat(sprintf('Gradient at initial theta (zeros): \n'))

## Gradient at initial theta (zeros):

cat(sprintf(' %f \n', grd))

## -0.100000
## -12.009217
## -11.262842

```

Part 3: Optimizing using optim

```

optimRes <- optim(par = initial_theta, fn = costFunction(X,y), gr = grad(X,y),
  method="BFGS", control = list(maxit = 400))

```

```

theta <- optimRes$par
cost <- optimRes$value

# Print theta to screen
cat(sprintf('Cost at theta found by optim: %f\n', cost))

## Cost at theta found by optim: 0.203498

cat(sprintf('theta: \n'))

## theta:

cat(sprintf(' %f \n', theta))

## -25.088453
## 0.205649
## 0.200882

mapFeature <- function(X1, X2) {
  # MAPFEATURE Feature mapping function to polynomial features
  #
  # MAPFEATURE(X1, X2) maps the two input features
  # to quadratic features used in the regularization exercise.
  #
  # Returns a new feature array with more features, comprising of
  # X1, X2, X1^2, X2^2, X1*X2, X1*X2^2, etc..
  #
  # Inputs X1, X2 must be the same size
  #

  degree <- 6

  out <- matrix(1,length(X1),1)

  for (i in 1:degree)
    for (j in 0:i)
      out <- cbind(out, (X1 ^ (i - j)) * (X2 ^ j))

  out
}

plotDecisionBoundary <-
function (theta, X, y, axLables = c("Exam 1 score","Exam 2 score"), legLabels =
  c('Admitted', 'Not admitted')) {
  # PLOTDECISIONBOUNDARY Plots the data points X and y into a new figure with
  # the decision boundary defined by theta
  # PLOTDECISIONBOUNDARY(theta, X,y) plots the data points with + for the
  # positive examples and o for the negative examples. X is assumed to be
  # a either
  # 1) Mx3 matrix, where the first column is an all-ones column for the
  # intercept.
  # 2) MxN, N>3 matrix, where the first column is all-ones

  # Plot Data
  plotData(X[,2:3], y,axLables,legLabels)

```

```

if (dim(X)[2] <= 3)
{
  # Only need 2 points to define a line, so choose two end points
  plot_x <- cbind(min(X[,2] - 2), max(X[,2] + 2))
  # Calculate the decision boundary line
  plot_y <- -1 / theta[3] * (theta[2] * plot_x + theta[1])

  # Plot, and adjust axes for better viewing
  lines(plot_x, plot_y, col = "blue")
}
else
{
  # Here is the grid range
  u <- seq(-1,1.5, length.out = 50)
  v <- seq(-1,1.5, length.out = 50)


  z <- matrix(0, length(u), length(v))
  # Evaluate z <- theta*x over the grid
  for (i in 1:length(u))
    for (j in 1:length(v))
      z[i,j] <- mapFeature(u[i], v[j]) %*% theta

  # Notice you need to specify the range [0, 0]
  contour(
    u, v, z, xlab = 'Microchip Test 1', ylab = 'Microchip Test 2',
    levels = 0, lwd = 2, add = TRUE, drawlabels = FALSE, col = "green"
  )
  mtext(paste("lambda = ",lambda), 3)
}
}


# Plot Boundary
plotDecisionBoundary(theta, X, y)

```

$X[, 2]$



$X[, 1]$



Part 4: Predict and Accuracies

```
predict <- function(theta, X) {
  #PREDICT Predict whether the label is 0 or 1 using learned logistic
  #regression parameters theta
  # p <- PREDICT(theta, X) computes the predictions for X using a
  # threshold at 0.5 (i.e., if sigmoid(theta'*x) >= 0.5, predict 1)

  m <- dim(X)[1] # Number of training examples
  p <- rep(0,m)
  p[sigmoid(X %*% theta) >= 0.5] <- 1
  p
  # -----
}
```

```
prob <- sigmoid(t(c(1,45,85)) %*% theta)
cat(sprintf('For a student with scores 45 and 85, we predict an admission probability of\n %f\n', prob))
```

```
## For a student with scores 45 and 85, we predict an admission probability of
## 0.775686
```

```
# Compute accuracy on our training set
```

```
p <- predict(theta, X)
```

```
cat(sprintf('Train Accuracy: %f\n', mean(p == y) * 100))
```




```
## Train Accuracy: 89.000000
```


Regularized Logistic Regression

Load and Plot Data

```
data <- read.table('ex2data2.txt', sep = ',')
X <- data[,c(1,2)]; y <- data[,3]
X <- as.matrix(X)

plotData(X, y, axLables = c('Microchip Test 1', 'Microchip Test 2'),
         legLabels = c("y = 1", "y = 0"))
```





Regularized

```
# Add Polynomial Features

# Note that mapFeature also adds a column of ones for us, so the intercept
# term is handled
X <- mapFeature(X[,1], X[,2])

# Initialize fitting parameters
initial_theta <- rep(0,dim(X)[2])

# Set regularization parameter lambda to 1
lambda <- 1

costFunctionReg <- function(X, y, lambda) {
  #COSTFUNCTIONREG Compute cost for logistic regression with regularization
  # J <- COSTFUNCTIONREG(theta, X, y, lambda) computes the cost of using
  # theta as the parameter for regularized logistic regression
  function(theta) {
    # Initialize some useful values
    m <- length(y) # number of training examples

    J <- 0
    # calculate hypothesis function h(x)
    h <- sigmoid(X %*% theta)
```

```

    # excluded the first theta value
    theta1 <- c(0,theta[-1])
    p <- lambda * (t(theta1) %*% theta1) / (2 * m)
    J <- (t(-y) %*% log(h) - t(1 - y) %*% log(1 - h)) / m + p
    J
    # -----
  }
}

gradReg <- function (X, y, lambda) {
  #COSTFUNCTIONREG Compute gradient for logistic regression with regularization
  # J <- COSTFUNCTIONREG(theta, X, y, lambda) computes the
  # gradient of the cost w.r.t. to the parameters.
  function(theta) {
    # Initialize some useful values
    m <- length(y); # number of training examples

    # You need to return the following variables correctly
    grad <- rep(0,length(theta))

    # calculate hypothesis function
    h <- sigmoid(X %*% theta)
    # excluded the first theta value
    theta1 <- c(0,theta[-1])

    # calculate grads
    grad <- (t(X) %*% (h - y) + lambda * theta1) / m
    grad
    # -----
  }
}

# Compute and display initial cost and gradient for regularized logistic
# regression
cost <- costFunctionReg(X, y, lambda)(initial_theta)
grd <- gradReg(X,y, lambda)(initial_theta)
cat(sprintf('Cost at initial theta (zeros): %f\n', cost))

```

```
## Cost at initial theta (zeros): 0.693147
```

Part 2: Regularization and Accuracies

```

# Initialize fitting parameters
initial_theta <- rep(0,dim(X)[2])

# Set regularization parameter lambda to 1 (you should vary this)
lambda <- 0
#try with lambda in (1,0,100,155)
# Optimize
optimRes <- optim(par = initial_theta,
                  fn = costFunctionReg(X,y,lambda),
                  gr = gradReg(X,y,lambda),
                  method="BFGS",


```

```
control = list(maxit = 400))

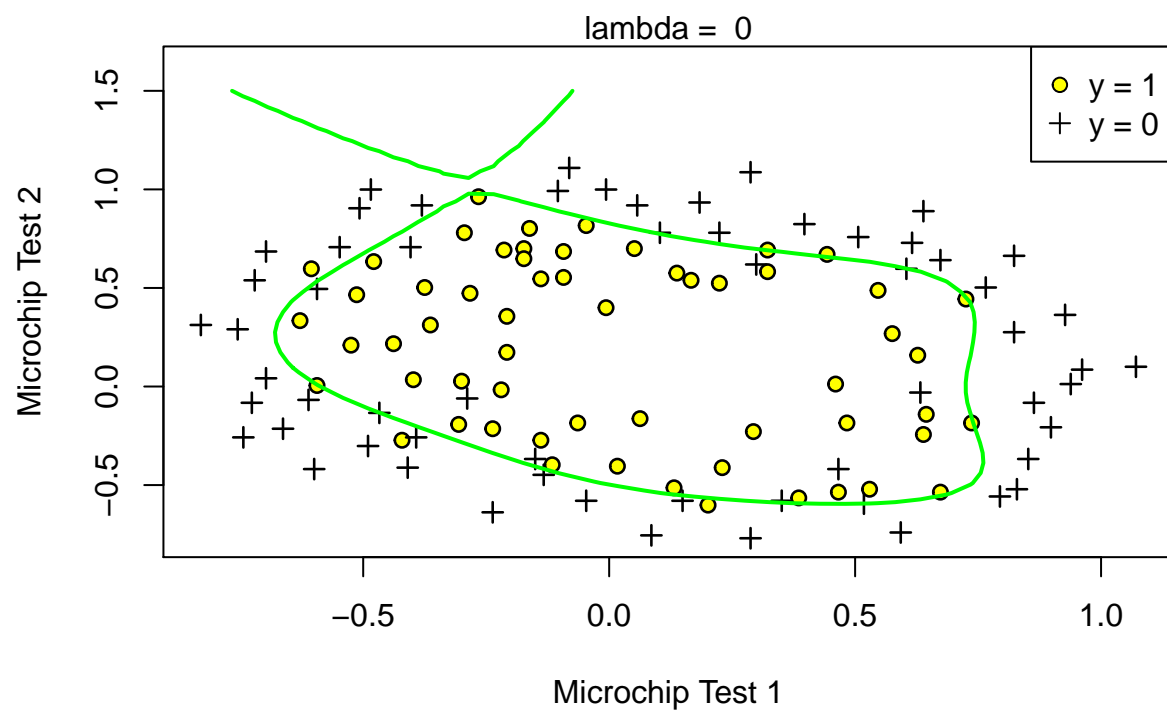
theta <- optimRes$par
J <- optimRes$value

# Plot Boundary
plotDecisionBoundary(theta, X, y, axLabels = c('Microchip Test 1', 'Microchip Test 2'),
  legLabels = c("y = 1", "y = 0"))
```

X[, 2]



X[, 1]



```
# Compute accuracy on our training set
p <- predict(theta, X)

cat(sprintf('Train Accuracy: %f\n', mean(p == y) * 100))
```

```
## Train Accuracy: 86.440678
```